



AMM: learned projection views and axis-wise interaction

a Expand each branch representation

After preceding MCA



b Mix axes, then aggregate



S: projection mixing

Across views within a branch

M: modality mixing

Across branches at each (s, d)

D: channel mixing

Within each (m, s) vector

M uses one learned 3-by-3 routing matrix shared across S and D.

S is a learned projection axis. Semantic roles of individual views are not established.